

F# Data: Making structured data first-class citizens

Tomas Petricek

University of Cambridge
tomas@tomaspetricek.net

Gustavo Guerra

Microsoft Corporation, London
gustavo@codebeside.org

Don Syme

Microsoft Research, Cambridge
dsyme@microsoft.com

Abstract

Most statically typed languages assume that programs run in a closed world, but this is not the case. Modern applications interact with external services and often access data in structured formats such as XML, CSV and JSON. Static type systems do not understand such external data sources and only make data access more cumbersome. Should we just give up and leave the messy world of external data to dynamic typing and runtime checks?

Of course, we should not give up! In this paper, we show how to integrate external data sources into the F# type system. As most real-world data on the web do not come with an explicit schema, we develop a type inference algorithm that infers a type from representative samples. Next, we use the type provider mechanism for integrating the inferred structures into the F# type system.

The resulting library significantly reduces the amount of code that developers need to write when accessing data. It also provides additional safety guarantees. Arguably, as much as possible if we abandon the incorrect closed world assumption.

Categories and Subject Descriptors D.3.3 [Programming Languages]: Language Constructs and Features

Keywords F#, Type Providers, JSON, XML, Data, Type Inference

1. Introduction

The key function of many modern applications is to connect multiple data sources or services and present the aggregated data in some form. Mobile applications for taking notes, searching train schedules or finding tomorrow's weather all communicate with one or more services over the internet.

Increasing number of such services provide REST-based endpoints that return data as CSV, XML or JSON. Despite numerous schematization efforts, the services typically do not come with schema. At best, the documentation provides a number of typical requests and sample responses.

For example, the OpenWeatherMap service provides an endpoint for getting current weather for a given city¹. The page documents the input URL parameters and shows one sample JSON response to illustrate the response structure. Using a standard library for working with JSON and HTTP, we might call the service and read the temperature as follows²:

```
let data = Http.Request("http://weather.org/?q=Prague")
match JsonValue.Parse(data) with
| Record(root) →
  match Map.find "main" root with
  | Record(main) →
    match Map.find "temp" main with
    | Number(num) → printfn "Lovely %f degrees!" num
    | _ → failwith "Incorrect format"
  | _ → failwith "Incorrect format"
| _ → failwith "Incorrect format"
```

The code assumes that the response has a particular format described in the documentation. The root node must be a record with a "main" field, which has to be another record containing a "temp" field with numerical value. When the format is incorrect, the data access simply fails with an exception.

Using the JSON type provider from the F# Data library, we can write code with exactly the same functionality in two lines:

```
type W = JsonProvider<"http://weather.org/?q=Prague">
printfn "Lovely %f degrees!" (W.GetSample().Main.Temp)
```

On the first line, `JsonProvider<"...">` invokes a type provider at compile-time with the URL as a sample. The type provider infers the structure of the response and provides a type with a `GetSample` method that returns a parsed JSON with nested properties `Main.Temp`, returning the temperature as a number.

In the rest of the paper, we give more detailed description of the mechanism and we also discuss theoretical safety properties of the approach. The key novel contributions of this paper are:

- We present type providers for XML, CSV and JSON (Section 2) that are available in the F# Data library and we also cover practical aspects of the implementation that contributed to their industrial adoption (Section 6).
- We describe a predictable type inference algorithm for structured data formats that underlies the type providers (Section 3). It is based on finding common supertype of a set of examples.
- We provide a minimal formal model (Section 4) and use it to prove a relativized notion of type safety for our type providers (Section 5). Although we focus on our specific case, the relativized safety property demonstrates a new way of thinking about ML-style type safety that is needed in the context of web.

Although the F# Data library has been widely adopted is one of the most downloaded F# libraries it has not yet been presented in an academic form. Additional documentation and source code can be found at <http://fsharp.github.io/FSharp.Data>.

¹ See "Current weather data": <http://openweathermap.org/current>

² We abbreviate the full URL: <http://api.openweathermap.org/data/2.5/weather?q=Prague&units=metric>

2. Structural type providers

We start with an informal overview that shows how F# Data type providers simplify working with JSON, XML and CSV. We also introduce the necessary aspects of F# 3.0 type providers along the way. The examples in this section illustrate a number of key properties of our type inference algorithm:

- Our mechanism is robust and predictable. This is important as the user directly works with the inferred types and should understand why a specific type was inferred from a given sample³.
- Our inference mechanism prefers records over unions. This better supports developer tooling – most F# editors provide code completion hints on “.” and so types with properties (records) are easier to use than types that require pattern matching.
- Finally, we handle a number of practical concerns that may appear overly complicated, but are important in the real world. This includes support for different numerical types, `null` values and missing data and also different ways of representing Booleans in CSV files.

2.1 Working with JSON documents

The JSON format used in the example in Section 1 is a popular format for data exchange on the web based on data structures used in JavaScript. The following is the definition of the `JsonValue` type used earlier to represent parsed JSON:

```
type JsonValue =  
    | Null  
    | Number of decimal  
    | String of string  
    | Boolean of bool  
    | Record of Map<string, JsonValue>  
    | Array of JsonValue[]
```

The OpenWeatherMap example in the introduction used only a (nested) record containing a numerical value. To demonstrate other aspects of the JSON type provider, we look at a more complex example that also involves `null` value and an array:

```
[ { "name": null, "age": 25 },  
  { "name": "Alexander", "age": 3.5 },  
  { "name": "Tomas" } ]
```

Say we want to print the names of people in the list with an age if it is available. As before, the standard approach would be to pattern match on the parsed `JsonValue`. The code would check that the top-level node is a `Array`, iterate over the elements checking that each is a `Record` with certain properties and throw an exception or skip values in incorrect format. The standard approach can be made nicer by defining helper functions. However, we still need to specify names of fields as strings, which is error prone and can not be statically checked.

Assuming the `people.json` file contains the above example and data is a string value that contains another data set in the same format, we can print names and ages as follows:

```
type People = JsonProvider<"people.json">  
  
let items = People.Parse(data)  
for item in items do  
    printf "%s " item.Name  
    Option.iter (printf "(%f)") item.Age
```

The code achieves the same simplicity as when using dynamically typed languages, but is statically type-checked. In contrast to the earlier example, we now use a local file `people.json` as a sample for the type inference, but then processes data from another source.

Type providers. The notation `JsonProvider<"people.json">` on the first line passes a *static parameter* to the type provider. Static parameters are resolved at compile-time, so the file name has to be a constant. The provider analyzes the sample and generates a type that we name `People`. In F# editors that use the F# Compiler Service [?], the type provider is also executed at development-time and so the same provided types are used in code completion.

The `JsonProvider` uses a type inference algorithm (Section 3) and infers the following types from the sample:

```
type Entity =  
    member Age : option<decimal>  
    member Name : option<string>  
  
type People =  
    member GetSample : unit → Entity[]  
    member Parse : string → Entity[]
```

The type `Entity` represents the person. The property `Name` is optional, because one of the sample records contains `null` as the name. The property `Age` is also optional, because the value is missing in one sample. The two sample ages are an integer `25` and a decimal `3.5` and so the common inferred type is `decimal`.

The type `People` provides two methods for reading data. `GetSample` returns the sample used for the inference and `Parse` parses a JSON string containing data in the same format as the sample. Since the sample JSON is a collection of records, both methods return an array of `Entity` values.

Error handling. In addition to the structure of the types, the type provider also specifies what code should be executed at run-time in place of `item.Name` and other operations. This is done through a *type erasure* mechanism demonstrated in Section 2.2. In this example, the runtime behaviour is the same as in the hand-written sample in Section 1 – a member access throws an exception if the document does not have the expected format.

Informally, the safety property discussed in Section 5 states that if the inputs are subtypes of some of the provided samples (*i.e.* the samples are representative), then no exceptions will occur. In other words, we cannot avoid all failures, but we can prevent some – for example, if the OpenWeatherMap changes the response format, the sample in Section 1 will not re-compile and the user knows that the code needs to be changed. What this means for the traditional ML type safety is discussed in Section 5.1.

The role of records. The sample code is easy to write thanks to the fact that most F# editors provide code completion when “.” is typed. The developer does not need to look at the sample JSON file to see what fields are available for a person. To support this scenario, our inference algorithm prefers records (we also treat XML elements and CSV rows as records).

In the above example, this is demonstrated by the fact that `Age` is marked as optional. An alternative is to provide two different record types (one with `Name` and other with `Name` and `Age`), but this would complicate the processing code.

2.2 Reading CSV files

In the CSV file format, the structure is a collection of rows (records) consisting of fields (with names specified by the first row). The inference needs to infer the types of fields. For example:

Ozone	Temp	Date	Autofilled
41,	67,	2012-05-01,	0
36.3,	72,	2012-05-02,	1
12.1,	74,	3 kveten,	0
17.5,	#N/A,	2012-05-04,	0

³ In particular, we do not use probabilistic methods where adding additional sample could change the shape of the type.

One difference between JSON and CSV formats is that in CSV, the literals have no data types. In JSON, strings are "quoted" and Booleans are `true` and `false`. This is not the case for CSV and so we need to infer not just the structure, but also the primitive types.

Assuming the sample is saved as `airdata.csv`, the following snippet prints all rows from another file that were not autofilled:

```
type AirCsv = CsvProvider<"airdata.csv">
let air = AirCsv.Parse(data)
for row in air.Rows do
    if not row.Autofilled then
        printf "%s: %d" row.Date row.Ozone
```

The type of the record (row) is, again, inferred from the sample. The Date column uses mixed formats and is inferred as string (although we support many date formats and "May 3" would work). More interestingly, we also infer Autofilled as Boolean, because the sample contains only 0 and 1 values and using those for Booleans is a common CSV convention. Finally, the fields Ozone and Temp have types decimal and option<int>.

Erasing type providers. At runtime, the type providers we describe use an erasure mechanism similar to Java Generics [?]. A type provider also generates code that is executed in place of the members of the generated types. In the above example, the compiled (and actually executed) code looks as follows:

```
let air = CsvFile.Load("airdata.csv", fun r →
    asDecimal r.[0], asIntOpt r.[1], r.[2], asBool r.[3])
for c1,_,c3,c4 in air.Rows do
    if not c4 then printf "%s: %d" c3 c1
```

The generated type `AirCsv` is erased to a type `CsvFile< α >`. The Load method takes the file name together with a function that turns a row represented as an array of strings to a typed representation – here, a four-element tuple `decimal*option<int>*string*bool` (this type is also used as a type argument of `CsvFile< α >`). The properties such as `Ozone` are then replaced with an accessor that reads the corresponding tuple element.

The type erasure mechanism is explained in our simplified formal model in Section 4. More details about the full implementation in F# can be found in the paper on F# type providers [?].

2.3 Processing XML data

The XML format is similar to JSON in that it uses nested elements rather than a flat CSV-like structure. In our inference algorithm, we use the notion of *records* for both JSON records and XML elements. However, unlike in JSON, the XML elements also have a name. An element can contain a collection of child elements. In case of XML, the collection can often be heterogeneous containing child nodes of different names. Consider the following (simplified) RSS news feed as an example:

```
<rss version="2.0"><channel>
  <title> BBC News - Europe </title>
  <item><title> Kurdish activists
    killed in Paris </title></item>
  <item><title> German MPs warn
    over UK EU exit </title></item>
</channel></rss>
```

The channel node contains a single title node (with the name of the feed) and multiple item nodes (individual news articles). In RSS, the title nodes always contain plain text, while item nodes contain title (and also url and other properties omitted here).

When extracting data from such XML documents, we often need to get nodes of a specific name (all items) or extract the content of a node (feed title). The XML type provider is optimized for this style of access and allows processing the feed as follows:

```
type RssFeed = XmlProvider<"rss-sample.xml">
let channel = RssFeed.Load("http://bbc.co.uk/...").Channel
printf "News from: %s " channel.Title
for item in channel.Items do
    printf " - %s " item.Title
```

In this example, the type of the channel value represents a heterogeneous collection that contains exactly one title node and zero or more item nodes.

Heterogeneous collections and unions. Our inference algorithm recognizes this as a *heterogeneous collection* (if the same structure appears in all samples) and generates a type with a member Title for accessing the value of the singleton title element together with a member Items that returns a collection of zero or more items. As a further simplification, if an element contains only a primitive value (such as title) it is represented as a value of primitive type (so channel.Title has a type string).

As an alternative, the inference could produce a collection of union type (with a case for item and a case for title). This alternative is useful for documents with complicated structure (such as XHTML files), but less useful for files with simpler structure such as web server responses, data stores and configuration files.

Moreover, our approach based on heterogeneous collections has a nice property that it generates types with the same public interface for the following two ways of representing data in XML:

```
<author name="Tomas" age="27" />
<author><name>Tomas</name><age>27</age></author>
```

In both cases, the provided type for author has two properties, typed string and int respectively. In Section 3, we discuss the simplified model (using union types). The inference based on heterogeneous collections is described later in Section 6.

2.4 Summary

The previous three sections demonstrated the F# Data type providers for JSON, CSV and XML. A reader interested in more examples is invited to look at documentation on the F# Data web site⁴.

It should be noted that we did not attempt to present the library using well-structured ideal samples, but instead used data sets that demonstrate typical issues that are frequent in real world inputs (missing data, inconsistent encoding of Booleans and heterogeneous structures).

From looking at just the previous three examples, these may appear arbitrary, but our experience suggests that they are the most common issues. The following JSON response with government debt information returned by the World Bank API [?] demonstrates all three problems:

```
[ { "page": 1, "pages": 5 },
  [ { "indicator": "GC.DOD.TOTL.GD.ZS",
    "date": "2012",
    "value": null },
    { "indicator": "GC.DOD.TOTL.GD.ZS",
    "date": "2010",
    "value": "35.1422970266502" } ] ]
```

First of all, the top-level element is a collection containing two values of different kind. The first is a record with meta-data about the current page and the second is an array with data. The JSON type provider infers this as heterogeneous collection with properties Record for the former and Array for the latter. Second, the value is `null` for some records. Third, numbers can be represented in JSON as numeric literals (without quotes), but here, they are returned as string literals instead⁵.

⁴ Available at <http://fsharp.github.io/FSharp.Data/>

⁵ This is often used to avoid non-standard numerical types in JavaScript.

3. Structural type inference

Our type inference algorithm for structured formats is based on a subtyping relation. When inferring the type of a document, we infer the most specific types of individual values (CSV rows, JSON or XML nodes) and find a common supertype of values in the sample.

In this section, we define the *structural type* σ , which is a type of structured data. Note that this type is distinct from the programming language types τ (the type generates the latter from the former). Next, we define the subtyping relation on structural types σ and describe the algorithm for finding a common supertype.

Note that the subtyping relation between structural types σ does not map to a subtyping relation between F# types τ . However, it does hold at runtime, meaning that an operation on a supertype can be performed on its subtypes. We make the type provider mechanism and runtime aspects precise in Section 4 before discussing safety in Section 5.

3.1 Structural types

The grammar below defines *structural type* σ . We distinguish between *non-nullable types* that always have a valid value (written as $\hat{\sigma}$) and *nullable types* that encompass missing and `null` values (written as σ). We write ν for record field names and ν_{opt} for record names, which can be empty:

$$\begin{aligned} \hat{\sigma} &= \nu_{\text{opt}} \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} \\ &\quad | \text{float} \mid \text{decimal} \mid \text{int} \mid \text{bit} \mid \text{bool} \mid \text{string} \\ \sigma &= \hat{\sigma} \mid \hat{\sigma} \text{ option} \mid [\sigma] \\ &\quad | \sigma_1 + \dots + \sigma_n \mid \top \mid \text{null} \end{aligned}$$

The non-nullable types include records (consisting of an optional name and zero or more fields with their types) and primitive types. Records arising from XML documents are always named, while records used by JSON and CSV providers are always unnamed.

Next, we have three numerical types – `int` for integers, `decimal` for small high-precision decimal numbers and `float` for floating-point numbers (these are related by sub-typing relation as discussed in the next section). Furthermore, the `bit` type is a type of values 0 and 1 and makes it possible to infer Boolean in the CSV processing example discussed in Section 2.2.

Any non-nullable type is also a nullable type, but it can be wrapped in the `option` constructor to explicitly permit the `null` value. These are typically mapped to the standard F# option type. A simple collection type $[\sigma]$ is also nullable and missing values or `null` are treated as empty collection. The type `null` is inhabited by the `null` value (using an overloaded but not ambiguous notation) and \top represents the top type.

Finally, a union type in our model implicitly permits the `null` value. This is because structural union types are *not* mapped to F# union types. Instead our encoding requires the user to handle the cases one-by-one and so they also need to handle the situation when none of the cases matches (as discussed in Section 6, this is the right choice for an F# type provider, but it would not necessarily be a good fit for other languages).

3.2 Subtyping relation

To provide a basic intuition, the subtyping relation between structured types is illustrated in Figure 1. We split the diagram in two parts. The upper part shows non-nullable types (with records and primitive types). The lower part shows nullable types with `null`, collections and optional values. We omit links between the two part, but any type $\hat{\sigma}$ is a subtype of $\hat{\sigma} \text{ option}$ (in the diagram, we abbreviate $\sigma \text{ option}$ as $\sigma?$).

The diagram shows only the basic structure and it does not explain relationships between records with the same name and between union types. This following definition makes this precise.

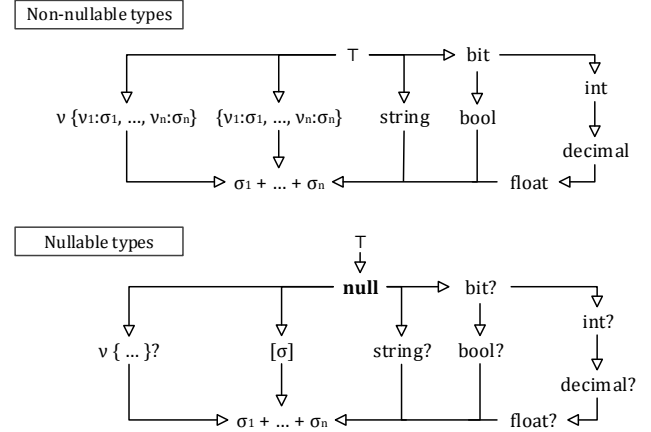


Figure 1. Subtype relation between structural types

Definition 1. We write $\sigma_1 :> \sigma_2$ to denote that σ_2 is a subtype of σ_1 . The subtyping relation is defined as a transitive reflexive closure of these rules:

Primitives, options, collections

$$\text{float} :> \text{decimal} :> \text{int} :> \text{bit} \quad \text{and} \quad \text{bool} :> \text{bit} \quad (\text{B1})$$

$$\sigma :> \text{null} \quad (\text{iff } \sigma \neq \hat{\sigma}) \quad (\text{B2})$$

$$\hat{\sigma} \text{ option} :> \hat{\sigma} \quad (\text{for all } \hat{\sigma}) \quad (\text{B3})$$

$$\sigma :> \top \quad (\text{for all } \sigma) \quad (\text{B4})$$

$$[\sigma_1] :> [\sigma_2] \quad (\text{if } \sigma_1 :> \sigma_2) \quad (\text{B5})$$

Union types

$$\sigma_1 + \dots + \sigma_n :> \sigma'_1 + \dots + \sigma'_n \quad (\text{iff } \forall i. \sigma_i :> \sigma'_i) \quad (\text{U1})$$

$$\sigma_1 + \dots + \sigma_n :> \sigma_1 + \dots + \sigma_m \quad (\text{where } m \leq n) \quad (\text{U2})$$

$$\sigma_1 + \dots + \sigma_n :> \sigma_{\pi(1)} + \dots + \sigma_{\pi(n)} \quad (\text{permutation } \pi) \quad (\text{U3})$$

$$\sigma_1 + \dots + \sigma_n :> \sigma_i \quad (\text{for } i \in 1 \dots n) \quad (\text{U4})$$

Record types

$$\nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} :> \nu \{ \nu_1 : \sigma'_1, \dots, \nu_n : \sigma'_n \} \quad (\sigma_i :> \sigma'_i) \quad (\text{R1})$$

$$\nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} :> \nu \{ \nu_1 : \sigma_1, \dots, \nu_m : \sigma_m \} \quad (m \geq n) \quad (\text{R2})$$

$$\begin{aligned} \nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} :> \\ \nu \{ \nu_{\pi(1)} : \sigma_{\pi(1)}, \dots, \nu_{\pi(m)} : \sigma_{\pi(m)} \} \end{aligned} \quad (\text{permutation } \pi) \quad (\text{R3})$$

$$\begin{aligned} \nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n, \nu_{n+1} : \text{null}, \dots, \nu_{n+m} : \text{null} \} :> \\ \nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} \end{aligned} \quad (\text{R4})$$

Here is a summary of the key aspects of the definition:

- For numeric types (B1), we infer a single most precise numeric type that can represent all values from a sample dataset, so we order types by the ranges they represent; `int` is a 32-bit integer, `decimal` has a range of about -1^{29} to 1^{29} and `float` is a double-precision floating-point.
- The `bit` type is a subtype of both `int` and `bool` (B1). Thus, a sample 0, 1 has a type `bit`; 0, 1, `true` is `bool` and 0, 1, 2 becomes `int` (for a sample 0, 1, 2, `true` we would need a union type).
- The `null` type is a subtype of all nullable types (B2), that is all σ types excluding non-nullable types $\hat{\sigma}$. Any non-nullable type is also a subtype of its optional version (B3).

(record-1)	$\frac{(\nu_i = \nu'_j \Leftrightarrow (i = j) \wedge (i \leq k)) \quad \forall i \in \{1..k\}.(\sigma_i \nabla \sigma'_i \vdash \sigma''_i)}{\nu \{ \nu_1 : \sigma_1, \dots, \nu_k : \sigma_k, \dots, \nu_n : \tau_n \} \nabla \nu \{ \nu'_1 : \sigma'_1, \dots, \nu'_k : \sigma'_k, \dots, \nu'_m : \tau'_m \} \vdash \nu \{ \nu_1 : \sigma''_1, \dots, \nu_k : \sigma''_k, \nu_{k+1} : [\sigma_{k+1}], \dots, \nu_n : [\sigma_n], \nu'_{k+1} : [\sigma'_{k+1}], \dots, \nu'_m : [\sigma'_m] \}}$	
(union-1)	$\frac{\exists i. \text{tagof}(\sigma_i) = \text{tagof}(\sigma) \quad \sigma \nabla \sigma_i \vdash \sigma'_i \quad \text{tagof}(\sigma) \neq \text{union}}{\sigma \nabla (\sigma_1 + \dots + \sigma_n) \vdash (\sigma_1 + \dots + [\sigma'_i] + \dots + \sigma_n)}$	$\frac{\nexists i. \text{tagof}(\sigma_i) = \text{tagof}(\sigma) \quad \text{tagof}(\sigma) \neq \text{union}}{\sigma \nabla (\sigma_1 + \dots + \sigma_n) \vdash (\sigma_1 + \dots + \sigma_n + [\sigma])}$
(union-2)	$\frac{\text{tagof}(\sigma_i) = \text{tagof}(\sigma'_i) \quad \sigma_i \nabla \sigma'_i \vdash \sigma''_i \quad (\forall i \in \{1..k\})}{(\sigma_1 + \dots + \sigma_k + \dots + \sigma_n) \nabla (\sigma'_1 + \dots + \sigma'_k + \dots + \sigma'_m) \vdash (\sigma''_1 + \dots + \sigma''_k + \sigma_{k+1} + \dots + \sigma_n + \sigma'_{k+1} + \dots + \sigma'_m)}$	
(union-3)	$\frac{(\text{no other rule applies})}{\sigma_1 \nabla \sigma_2 \vdash [\sigma_1] + [\sigma_2]}$	
(order)	$\frac{\nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} \nabla \sigma \vdash \sigma'}{\nu \{ \nu_{\pi(1)} : \sigma_{\pi(1)}, \dots, \nu_{\pi(n)} : \sigma_{\pi(n)} \} \nabla \sigma \vdash \sigma'}$	$\frac{\sigma_1 + \dots + \sigma_n \nabla \sigma \vdash \sigma'}{\sigma_{\pi(1)} + \dots + \sigma_{\pi(n)} \nabla \sigma \vdash \sigma'} \quad (\pi \text{ permutation})$
(list)	$\frac{\sigma_1 \nabla \sigma_2 \vdash \sigma}{[\sigma_1] \nabla [\sigma_2] \vdash [\sigma]}$	$\text{(prim)} \quad \frac{\sigma_1 >: \sigma_2}{\sigma_1 \nabla \sigma_2 \vdash \sigma_1} \quad \text{(when } \sigma_1, \sigma_2 \in \{\text{bit, int, decimal, float}\} \text{ or } \sigma_1, \sigma_2 \in \{\text{bit, bool}\})$
(sym)	$\frac{\sigma_1 \nabla \sigma_2 \vdash \sigma}{\sigma_2 \nabla \sigma_1 \vdash \sigma}$	$\text{(refl)} \quad \sigma \nabla \sigma \vdash \sigma \quad \text{(null-1)} \quad \sigma \nabla \text{null} \vdash \sigma \quad (\sigma >: \text{null})$
	$\text{(top)} \quad \top \nabla \sigma \vdash \sigma$	$\text{(null-2)} \quad \sigma \nabla \text{null} \vdash \sigma \text{ option} \quad (\sigma \not>: \text{null})$

Figure 2. Inference judgements that define the common supertype relation

- There is a top type (B4), but no bottom type. It is possible to find common supertype of any two types, because a union type $\tau_1 + \dots + \tau_n$ is a supertype of all its components (U4).
- The remaining rules for unions are standard. Subtype can have fewer alternatives (U1), elements can be reordered (U2) and subtyping is covariant (U3). Similarly, the collection type is also covariant (B5).
- As usual, the subtyping on records is covariant (R1), subtype can have additional fields (R2) and fields can be reordered (R3). The interesting rule is the last one (R4). Together with covariance, it states that a subtype can omit some fields, provided that their types are nullable.

The rule that allows subtype to have fewer record elements (R4) is particularly important. It allows us to prefer records in some cases. For example, given two samples $\{\text{name} : \text{string}\}$ and $\{\text{name} : \text{string}, \text{age} : \text{int}\}$, we can find a common supertype $\{\text{name} : \text{string}, \text{age} : \text{int option}\}$ which is also a record. For usability reasons, we prefer this to another common supertype $\{\text{name} : \text{string}\} + \{\text{name} : \text{string}, \text{age} : \text{int}\}$. The next section describes precisely how our inference algorithm works.

It is worth noting that some of the aspects (such as 3 different numeric types and handling of the `null` type) are specific to F#. Similar problems will appear with most real-world languages and systems. JVM and OCaml all have multiple numeric types, but they would handle `null` differently.

We could also be more or less strict about missing data – we choose to handle missing values silently when we can (`null` becomes an empty collection), but we are explicit in other cases (we infer `string option` as a hint to the user rather than treating `null` as an empty string). These choices are based on experience with using F# Data in practice.

3.3 Common supertype relation

As mentioned earlier, the structured type inference relies on a finding common supertype. However, the partially ordered set of types does not have a unique greatest lower bound. For example, record types $\{a : \text{int}\}$ and $\{b : \text{bool}\}$ have common supertypes

$\{?a : \text{int}, ?b : \text{bool}\}$ and $\{a : \text{int}\} + \{b : \text{bool}\}$ which are unrelated. Our inference algorithm always prefers records over unions (Theorem 3). The definition of the *common supertype* relation uses a number of auxiliary definitions that are explained in the discussion that follows.

Definition 2. A common supertype of types σ_1 and σ_2 is a type σ , written $\sigma_1 \nabla \sigma_2 \vdash \sigma$, obtained according to the inference rules in Figure 2.

When finding a common supertype of two records (*record-1*), we return a record type that has the union of fields of the two arguments. We assume that the names of the first k fields are the same and the remaining fields have different names. To get a record with the right order of elements, we provide the (*order*) rule. The types of shared fields become common supertypes of their respective types (recursively). Fields that are present in only one record are marked as optional using the following helper definition:

$$\begin{aligned} [\hat{\sigma}] &= \hat{\sigma} \text{ option} & (\text{non-nullable types}) \\ [\sigma] &= \sigma & (\text{otherwise}) \end{aligned}$$

Finding a common supertype of when one type is a union is more complicated. Our definition aims to use a flat structure (avoid nesting unions) and limit the number of cases. This is done by grouping types that have a common supertype which is not a union type. For example, rather than inferring $\text{int} + (\text{bool} + \text{decimal})$, our algorithm will find the common supertype of int and decimal (which is decimal) and it will produce just $\text{decimal} + \text{bool}$.

To identify which types have a common supertype which is not a union, we group the types by a *tag*, which is defined as:

$$\begin{aligned} \text{tag} &= \text{string} \mid \text{bool} \mid \text{number} \mid \text{union} \mid \text{collection} \\ &\quad \mid \text{rec-anon} \mid \text{rec-named } \nu \end{aligned}$$

The tag of a type is obtained using a function $\text{tagof}(-) : \sigma \rightarrow \text{tag}$:

$$\begin{aligned} \text{tagof}(\text{string}) &= \text{string} \\ \text{tagof}(\text{bool}) &= \text{bool} \\ \text{tagof}(\text{decimal}) &= \text{tagof}(\text{float}) = \text{number} \\ \text{tagof}(\text{int}) &= \text{tagof}(\text{bit}) = \text{number} \\ \text{tagof}([\sigma]) &= \text{collection} \end{aligned}$$

$\text{asInt}(i) \rightsquigarrow i$	$\text{asDec}(i) \rightsquigarrow d \quad (d = i)$	$\text{asFloat}(i) \rightsquigarrow f \quad (f = i)$	$\text{asStr}(i) \rightsquigarrow t \quad (t \text{ represents } i)$
$\text{asInt}(d) \rightsquigarrow i \quad (i = \lfloor d \rfloor)$	$\text{asDec}(d) \rightsquigarrow d$	$\text{asFloat}(d) \rightsquigarrow f \quad (f = d)$	$\text{asStr}(d) \rightsquigarrow t \quad (t \text{ represents } d)$
$\text{asInt}(f) \rightsquigarrow i \quad (i = \lfloor f \rfloor)$	$\text{asDec}(f) \rightsquigarrow d \quad (d = \lfloor f \rfloor)$	$\text{asFloat}(f) \rightsquigarrow f$	$\text{asStr}(f) \rightsquigarrow t \quad (t \text{ represents } f)$
$\text{asInt}(t) \rightsquigarrow i \quad (t \text{ represents } i)$	$\text{asDec}(t) \rightsquigarrow d \quad (t \text{ represents } d)$	$\text{asFloat}(t) \rightsquigarrow f \quad (t \text{ represents } f)$	$\text{asStr}(t) \rightsquigarrow t$
$\text{asInt}(\text{true}) \rightsquigarrow 1$	$\text{asDec}(\text{true}) \rightsquigarrow 1.0$	$\text{asFloat}(\text{true}) \rightsquigarrow 1.0$	$\text{asStr}(\text{true}) \rightsquigarrow \text{"true"}$
$\text{asInt}(\text{false}) \rightsquigarrow 0$	$\text{asDec}(\text{false}) \rightsquigarrow 0.0$	$\text{asFloat}(\text{false}) \rightsquigarrow 0.0$	$\text{asStr}(\text{false}) \rightsquigarrow \text{"false"}$
$\text{asBool}(1) \rightsquigarrow \text{true}$	$\text{asBool}(0) \rightsquigarrow \text{false}$		
$\text{asBool}(1.0) \rightsquigarrow \text{true}$	$\text{asBool}(0.0) \rightsquigarrow \text{false}$		
$\text{asBool}(\text{"true"}) \rightsquigarrow \text{true}$	$\text{asBool}(\text{"false"}) \rightsquigarrow \text{false}$		
$\text{asBool}(\text{"1"}) \rightsquigarrow \text{true}$	$\text{asBool}(\text{"0"}) \rightsquigarrow \text{false}$		
$\text{isNull}(\text{null}) \rightsquigarrow \text{true}$	$\text{asBool}(b) \rightsquigarrow b$		
$\text{isNull}(_) \rightsquigarrow \text{false}$	$\text{getChildren}(\text{null}) \rightsquigarrow []$	$\text{getField}(\nu, \nu_i, \nu \{ \dots, \nu_i = s_i, \dots \}) \rightsquigarrow s_i$	
$\text{isTag}(\text{string}, t) \rightsquigarrow \text{true}$		$\text{getField}(\perp, \nu_i, \{ \dots, \nu_i = s_i, \dots \}) \rightsquigarrow s_i$	
$\text{isTag}(\text{bool}, v) \rightsquigarrow \text{true} \quad (\text{when } v \in 0, 1, 0.0, 1.0, \text{true}, \text{false})$		$\text{getField}(\nu, \nu', \nu \{ \dots, \nu_i = s_i, \dots \}) \rightsquigarrow \text{null} \quad (\nexists i. \nu_i = \nu')$	
$\text{isTag}(\text{number}, v) \rightsquigarrow \text{true} \quad (\text{when } v = i, v = d, v = f)$		$\text{getField}(\perp, \nu', \{ \dots, \nu_i = s_i, \dots \}) \rightsquigarrow \text{null} \quad (\nexists i. \nu_i = \nu')$	
$\text{isTag}(\text{collection}, [s_1; \dots; s_n]) \rightsquigarrow \text{true}$		$\text{getChildren}([s_1; \dots; s_n]) \rightsquigarrow [s_1; \dots; s_n]$	
		$\text{isTag}(\text{collection}, \text{null}) \rightsquigarrow \text{true}$	
		$\text{isTag}(\text{rec-anon}, \{ \nu_1 \mapsto s_1, \dots, \nu_n \mapsto s_n \}) \rightsquigarrow \text{true}$	
		$\text{isTag}(\text{rec-named } \nu, \nu \{ \nu_1 \mapsto s_1, \dots, \nu_n \mapsto s_n \}) \rightsquigarrow \text{true}$	
		$\text{isTag}(_, _) \rightsquigarrow \text{false}$	

Figure 3. Reduction rules for conversion functions

$$\begin{aligned}
&\text{tagof}(\sigma + \dots + \sigma) = \text{union} \\
&\text{tagof}(\{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \}) = \text{rec-anon} \\
&\text{tagof}(\nu \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \}) = \text{rec-named } \nu \\
&\text{tagof}(\hat{\sigma} \text{ option}) = \text{tagof}(\hat{\sigma})
\end{aligned}$$

The function is undefined for the \top and null types, but this is not a problem because these types never need to be used as arguments in Figure 2. The \top type is always eliminated using the (*top*) rule and the null type is eliminated using either (*null-1*) or (*null-2*).

The handling of unions is specified using three rules. When adding a non-union type to a union (*union-1*), the union may or may not contain a case with the same tag as the new type. If it does, the new type is combined with the existing one, otherwise a new case is added. When combining two unions (*union-2*), we group the cases that have a shared tags. Finally, the last rule (*union-3*) covers the case when we are combining two non-union types. As discussed earlier, union implicitly permits null values and so we use the following auxiliary function which makes nullable types non-nullable (when possible) to simplify the type:

$$\begin{aligned}
\lfloor \hat{\sigma} \text{ option} \rfloor &= \hat{\sigma} \quad (\text{option}) \\
\lfloor \sigma \rfloor &= \sigma \quad (\text{otherwise})
\end{aligned}$$

The remaining rules are straightforward. For collections, we find the common supertype of its elements (*list*); for compatible primitive types, we choose one of the two (*prim*) and a common supertype with null is either the type itself (*null-1*) or an option (*null-2*).

Properties. The common supertype relation finds a single supertype (among multiple candidates). The following theorems specify that the found type is uniquely defined (up to reordering of fields in records and union cases) and is, indeed, a common supertype.

Theorem 1 (Uniqueness). *If $\sigma_1 \nabla \sigma_2 \vdash \sigma$ and $\sigma_1 \nabla \sigma_2 \vdash \sigma'$ then it holds that $\sigma \succ \sigma'$ and $\sigma' \succ \sigma$.*

Proof. The assumptions and shapes of the types to be unified in the rules in Figure 2 are disjoint, with the exception of (*order*), (*sym*) and (*refl*). These rules always produce types that are subtypes of each other. This property is preserved by rules that use the common supertype relation recursively. \square

Theorem 2 (Common supertype). *If $\sigma_1 \nabla \sigma_2 \vdash \sigma$ then it holds that $\sigma \succ \sigma_1$ and $\sigma \succ \sigma_2$.*

Proof. By induction over the common supertype derivation \vdash . \square

We stated earlier that the common supertype relation minimises the use of union types (by preferring common numeric types or common record types when possible). This property can be stated and proved formally:

Theorem 3. *Given σ_1, σ_2 if there is $s \sigma$ such that $\sigma \succ \sigma_1$ and $\sigma \succ \sigma_2$ and σ is not a union type, then $\sigma_1 \nabla \sigma_2 \vdash \sigma'$ and σ' is not a union type.*

Proof. The only rule that introduces an union type is (*union-3*), which is used only when no other inference rule can be applied. \square

4. Formalising type providers

In this section, we build the necessary theoretical framework for proving the relativised type safety property of F# Data type providers in Section 5. We start by discussing the runtime representation of structured values and runtime conversions (Section 4.1). Then we embed these into a formal model of an F# subset that is necessary for discussing type providers (Section 4.2) and we describe how F# Data type providers turn an inferred structural type into actual F# code (Section 4.3).

4.1 Structural values and conversions

In the model presented here, we represent JSON, XML and CSV documents using the same *structural value*⁶. Structural values are first-order and can be one of the primitive values (Boolean, string, integer, decimal and float), a missing value (null), a collection and a record (named or unnamed):

$$\begin{aligned}
s &= i \mid d \mid f \mid t \mid \text{true} \mid \text{false} \mid \text{null} \\
&\mid [s_1; \dots; s_n] \\
&\mid \{ \nu_1 \mapsto s_1, \dots, \nu_n \mapsto s_n \} \\
&\mid \nu \{ \nu_1 \mapsto s_1, \dots, \nu_n \mapsto s_n \}
\end{aligned}$$

The first few cases represent primitive values (i for integers, d for decimals, f for floating point numbers and t for strings). A collection is written as a list of values in square brackets, separated by semicolons. A record can optionally start with its name ν , followed by a sequence of field assignments $\nu_i \mapsto s_i$ written in curly brackets.

⁶ Here, we diverge from the actual implementation in F# Data which uses a different implementation for each format (reusing existing libraries).

As suggested by subtyping, our system permits certain runtime conversions (e.g. a numerical value 1 can be treated as `true`). These are captured by the following operations:

```

op = asInt(s) | asDec(s) | asFloat(s) | asBool(s)
    | asStr(s) | getChildren(s) | getField( $\nu_{\text{opt}}, \nu, s$ )
    | isNull(s) | isTag(t, s)

```

The reduction rules for these operations are defined in Figure 3. It is important to note that the relativised property we prove about F# Data specifies *sufficient*, but not *necessary* conditions. That is, a type provided based on a sample is guaranteed to work correctly on certain inputs, but may also handle several additional inputs.

The conversion functions follow this approach. They perform conversions required by the subtyping relation (e.g. integer is converted to decimal or a float) but also attempt additional conversions.

- **Widening operations.** The widening operations are required by the subtyping. These include converting to a wider numerical type (integer to decimal and float, decimal to float, converting 0 and 1 to Boolean).
- **Conversions.** Additional conversions are not strictly required, but prove useful in practice. These include converting to a smaller numerical type (such as float to integer) written as $i = \lfloor f \rfloor$, which is not defined in case of arithmetic overflow; conversion to and from string values written as “ t represents i ”, which parses or formats a non-string value and also trating of decimals 0.0 and 1.0 as Booleans.
- **Structural operations.** Aside from primitives, the `getChildren` operation returns treats `null` as an empty collection. The `getField` function is defined separately for named and unnamed records (we write \perp for a name of an unnamed record). Accessing a field of a named record checks that the expected record name matches the actual.
- **Helper functions.** The semantics also defines two helper functions. The `isNull` function is used to check whether value is `null` and `isTag` is used to test whether a value can be treated as a value of type associated with the specified tag (as defined in Section 3.3). The last line defines a “catch all” pattern, which returns `false` for all remaining cases.

4.2 Featherweight F#

The semantics fragment introduced in the previous section discusses values and operations that are used by F# Data. This section adds a minimal subset of the standard F# language.

We focus on the object system (which is what F# type providers produce) and so the following combines standard ML [?] features with Featherweight Java [?] classes. However, we only need classes with properties and without inheritance. A class has a single implicit constructor and the declaration closes over constructor parameters. To avoid including all of ML, we only select constructs for working with options and lists that we need later.

```

 $\tau = \text{int} \mid \text{decimal} \mid \text{float} \mid \text{bool} \mid \text{string}$ 
     $\mid C \mid \text{StructVal} \mid \text{list}(\tau) \mid \text{option}(\tau)$ 

```

```

 $L = \text{type } C(\overline{x} : \overline{\tau}) = \overline{M}$ 
 $M = \text{member } N : \tau = e$ 

```

```

 $v = s \mid \text{None} \mid \text{Some}(v) \mid \text{new } C(\overline{v}) \mid [v_1; \dots; v_n]$ 
 $e = op \mid e.N \mid \text{new } C(\overline{e}) \mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3$ 
     $\mid \text{Some}(e) \mid \text{match } e \text{ with } \text{Some}(x) \rightarrow e_1 \mid \text{None} \rightarrow e_2$ 
     $\mid \text{match } e \text{ with } [x_1; \dots; x_n] \rightarrow e_1 \mid \_ \rightarrow e_2$ 
     $\mid [e_1; \dots; e_n] \mid \text{List.map } (\lambda x \rightarrow e_1) e_2$ 

```

In the above definition, τ denotes types including a type of all structured values s called `StructVal`. A class definition L consists of

```

(member)  $\frac{\text{type } C(\overline{x} : \overline{\tau}) = \dots \text{member } N_i : \tau_i = e_i \dots}{(\text{new } C(\overline{v})).N_i \rightsquigarrow e_i[\overline{x} \leftarrow \overline{v}]}$ 

(ctx)  $E[e] \rightsquigarrow E[e'] \quad (\text{when } e \rightsquigarrow e')$ 

(cond1)  $\text{if true then } e_1 \text{ else } e_2 \rightsquigarrow e_1$ 

(cond2)  $\text{if false then } e_1 \text{ else } e_2 \rightsquigarrow e_2$ 

(match1)  $\text{match None with } \text{Some}(x) \rightarrow e_1 \mid \text{None} \rightarrow e_2 \rightsquigarrow e_2$ 

(match2)  $\text{match Some}(v) \text{ with } \text{Some}(x) \rightarrow e_1 \mid \text{None} \rightarrow e_2 \rightsquigarrow e_1[x \leftarrow v]$ 

(match3)  $\text{match } [v_1; \dots; v_n] \text{ with } [x_1; \dots; x_m] \rightarrow e_1 \mid \_ \rightarrow e_2 \rightsquigarrow e_2 \quad (m \neq n)$ 

(match4)  $\text{match } [v_1; \dots; v_n] \text{ with } [x_1; \dots; x_n] \rightarrow e_1 \mid \_ \rightarrow e_2 \rightsquigarrow e_1[\overline{x} \leftarrow \overline{v}]$ 

(map)  $\text{List.map } (\lambda x.e) [v_1; \dots] \rightsquigarrow [e[x \leftarrow v_1]; \dots]$ 

```

Figure 4. Featherweight F# – Remaining reduction rules

a constructor and zero or more members. Values v include previously defined structural values s and values for the option and list type; finally expressions e include previously defined operations op , class construction, member access, conditionals and expressions for working with option values and lists. We include `List.map` as a special construct to avoid making the language too complex.

Next, we define the reduction relation and (a fragment of) type checking for Featherweight F#. The language presented here is intentionally incomplete. We only define parts needed to prove the relativized soundness property in Section 5.

Reduction. The reduction relation is of the form $e \rightsquigarrow e'$. We also write $e \rightsquigarrow^* e'$ to denote the reflexive and transitive closure of \rightsquigarrow . The reduction rules for primitive functions op were discussed earlier in Figure 3. The rules for other expressions defined above are provided in Figure 4.

The *(ctx)* rule performs a reduction inside a sub-expression specified by an evaluation context. This models the eager evaluation order of F#. An evaluation context E is defined as:

```

 $E = f(E) \mid E.N \mid \text{new } C(\overline{v}, E, \overline{e}) \mid \text{if } E \text{ then } e_1 \text{ else } e_2$ 
     $\mid \text{Some}(E) \mid \text{match } E \text{ with } \text{Some}(x) \rightarrow e_1 \mid \text{None} \rightarrow e_2$ 
     $\mid [\overline{v}; E; \overline{e}] \mid \text{match } E \text{ with } [x_1; \dots; x_n] \rightarrow e_1 \mid \_ \rightarrow e_2$ 
     $\mid \text{map } (\lambda x \rightarrow e) E \mid [\_]$ 

```

```

 $f \in \{ \text{asInt}, \text{asDec}, \text{asFloat}, \text{asBool}, \text{asStr},$ 
     $\text{getField}, \text{getChildren}, \text{isNull}, \text{isTag} \}$ 

```

The evaluation first reduces arguments of functions and the evaluation proceeds from left to right as denoted by $\overline{v}, E, \overline{e}$ in constructor arguments or $\overline{v}; E; \overline{e}$ in list initialization.

We write $e[\overline{x} \leftarrow \overline{v}]$ for the result of replacing variables \overline{x} by values \overline{v} in expression. The *(member)* rule reduces a member access using a class definition in the assumption to obtain the body of the expression. Finally, the remaining six rules provide standard reductions for conditionals and pattern matching.

The language is sufficient for our purpose, but simple. It does not provide recursion and so all expressions reduce to a value in a finite number of steps or get stuck due to an error condition. An error condition can be a wrong argument passed to conditional, pattern matching or one of the conversion functions from Figure 3.

$$\begin{array}{c}
\frac{}{\Gamma \vdash s : \text{StructVal}} \quad \frac{}{\Gamma \vdash n : \text{int}} \\
\\
\frac{\Gamma \vdash e : C \quad \text{type } C(\overline{x} : \overline{\tau}) = .. \text{ member } N_i : \tau_i = e_i ..}{\Gamma \vdash e.N_i : \tau_i} \\
\\
\frac{\Gamma \vdash e_i : \tau_i \quad \text{type } C(x_1 : \tau_1, \dots, x_n : \tau_n) = \dots}{\Gamma \vdash \text{new } C(e_1, \dots, e_n) : C}
\end{array}$$

Figure 5. Featherweight F# – Fragment of type checking

Type checking. Selected type checking rules for Featherweight F# are shown in Figure 5. The rules demonstrate the key difference between our language and standard ML or Featherweight Java:

- All structured values s have a type `StructVal`. Some of those have other types (primitive integers, decimals, floats, strings and Booleans as illustrated by the rule for n). For other values, `StructVal` is the only type – this includes records and `null`.
- A list containing other structural values $[s_1; \dots; s_n]$ has a type `StructVal`, but can also have the list $\langle \tau \rangle$ type. Conversely, lists that contain non-structural values like objects or options are not of type `StructVal`.
- The two rules for checking class construction and member access are similar to corresponding rules of Featherweight Java.

An important part of Featherweight Java that is omitted here is the checking of type declarations (ensuring the bodies of members are well-typed). We omit this, because it is not needed for the relativized safety theorem, which considers only classes generated by our type provider (which are correct by construction).

4.3 Type providers

So far, we defined the type inference algorithm for structured data formats which produces a structured type σ based on one or more sample documents (Section 3). In Section 4.2, we defined a simplified model of F# language and evaluation, which also included conversions that model the F# Data library runtime (Section 4.3). In this section, we define how the type providers work, which links the two aspects discussed so far.

The type providers for XML, CSV and JSON all work in the same way. They take sample documents and infer a common supertype σ from the samples. From σ , they then generate F# types that are then exposed to the programmer⁷.

Type provider mapping. When generating types, the type provider produces an F# type τ , expression that wraps a structured document value (of type `StructVal`) as a value of type τ and a collection of class definitions. We express it using the following mapping:

$$\llbracket - \rrbracket_e : \sigma \rightarrow \tau \times e' \times L$$

The semantics is parameterized by an expression e , which represents code to obtain a structured value that is being wrapped. Then, given an inferred structured type σ , the type provider produces an F# type τ , an expression e' which constructs a value of type τ using e and also a set of class definitions L .

Figure 6 shows the rules that define $\llbracket - \rrbracket_e$. Primitive types are all handled by a single rule. For a given structural type, it returns the corresponding F# type – the types are mapped directly with the exception of `bit`, which becomes an F# `bool`. The generated code calls an appropriate conversion function from Figure 3 on the input.

Handling of records is more interesting. We generate a new class C that takes a structured value as constructor parameter. For each record field, we generate a new member with the same name as the

field⁸. The body of the member calls `getField` and then passes this expression to $\llbracket \sigma_i \rrbracket$ which adds additional wrapping that maps the field (structured value of type σ_i) into an F# type τ_i . The returned expression creates a new instance of C and the mapping returns the class C together with all recursively generated definitions.

A collection type becomes an F# list $\langle \tau \rangle$. The returned expression calls `getChildren` (which turns `null` values into empty lists) and then uses `List.map` to convert all nested values to an F# type τ . The handling of option type is similar – it checks if the original value is `null` and if no, it wraps the recursively generated conversion expression e' in the `Some` constructor.

As discussed earlier, union types are also generated as classes with properties. Given a union type $\sigma_1 + \dots + \sigma_n$, we get corresponding F# types τ_i and generate n members of type `option` $\langle \tau_i \rangle$. Each member return `Some` when the value is not `null` and has the right structure (checked by `isTag`). The type inference algorithm also guarantees that there is only one case for each type tag (Section 3.3) and there are no nested union types. Thus, checking for a tag is sufficient and we can also use the tag to identify the name of the generated member (using the `nameof` function).

Example 1. To illustrate how the type provision mechanism works, we now consider two simple examples. First, assume that the inferred type is a record with two fields (one optional) such as `Person { Age : int option, Name : string }`. Applying the rules from Figure 6 produces the following class:

```

type Person(v : StructVal) =
  member Age : option<int> =
    if isNull(getField(Person, Age, v)) then None
    else Some(asInt(getField(Person, Age, v)))
  member Name : string =
    asStr(getField(Person, Name, v))

```

The body of the `Age` member is produced by the case for optional types applied to an expression `getField(Person, Age, v)`. If the returned field is not `null`, then the member calls `asInt` (produced by the case for primitive types) and wraps the result in the `Some` constructor. Note that `getField` is defined even when the field does not exist, but returns `null`. This lets us treat missing fields as optional fields. The `Name` member is similar, but does not perform any checks.

For completeness, a type corresponding to the record is `Person` and given a structural document s , we create a `Person` value by calling `new Person(s)`.

Example 2. The second example illustrates the remaining parts of the type provision, including collections and union types. Teusing the `Person` type from the previous example, consider a collection `[Person+string]`, which contains a mix of `Person` and string values.

```

type PersonOrString(v : StructVal) =
  member Person : option<Person> =
    if isNull(v) then None else
    if isTag(rec-named Person, v) then
      Some(new Person(v)) else None
  member String : option<string> =
    if isNull(v) then None else
    if isTag(string, v) then
      Some(asStr(v)) else None

```

⁷ The actual implementation provides *erased types* as described in [?].

Here, we treat the code as actually generated. This is an acceptable simplification, because F# Data type providers do not rely on laziness that is available through *erased types*.

⁸ The actual F# Data implementation also capitalizes the names.

$\llbracket \sigma_p \rrbracket_e = \tau_p, op(e), \emptyset \quad \text{where}$ $\sigma_p, \tau_p, op \in \{ (\text{bit}, \text{bool}, \text{asBool}), (\text{bool}, \text{bool}, \text{asBool}), (\text{int}, \text{int}, \text{asInt}), (\text{decimal}, \text{decimal}, \text{asDec}), (\text{float}, \text{float}, \text{asFloat}), (\text{string}, \text{string}, \text{asStr}) \}$	$\llbracket \hat{\sigma} \text{ option} \rrbracket_e =$ $\text{option} \langle \tau \rangle, \text{if isNull } e \text{ then None else Some}(e'), L$ $\text{where } \tau, e', L = \llbracket \hat{\sigma} \rrbracket_e$
$\llbracket \nu_{\text{opt}} \{ \nu_1 : \sigma_1, \dots, \nu_n : \sigma_n \} \rrbracket_e =$ $C, \text{new } C(e), L_1 \cup \dots \cup L_n \cup \{L\} \quad \text{where}$ $C \text{ is a fresh class name}$ $L = \text{type } C(v : \text{StructVal}) = M_1 \dots M_n$ $M_i = \text{member } \nu_i : \tau_i = e_i$ $\tau_i, e_i, L_i = \llbracket \sigma_i \rrbracket_{e'}, \quad e' = \text{getField}(\nu_{\text{opt}}, \nu_i, e)$	$\llbracket \sigma_1 + \dots + \sigma_n \rrbracket_e =$ $C, \text{new } C(e), L_1 \cup \dots \cup L_n \cup \{L\} \quad \text{where}$ $C \text{ is a fresh class name}$ $L = \text{type } C(v : \text{StructVal}) = M_1 \dots M_n$ $M_i = \text{member } \nu_i : \text{option} \langle \tau_i \rangle =$ $\text{if isNull}(v) \text{ then None else}$ $\text{if isTag}(t_i, v) \text{ then Some}(e_i) \text{ else None}$ $\tau_i, e_i, L_i = \llbracket \sigma_i \rrbracket_e, \quad t_i = \text{tagof}(\sigma_i), \quad \nu_i = \text{nameof}(t_i)$
$\llbracket [\sigma] \rrbracket_e = \text{list} \langle \tau \rangle, \text{List.map } (\lambda x \rightarrow e') (\text{getChildren}(e)), L$ $\text{where } \tau, e', L = \llbracket \hat{\sigma} \rrbracket_x$	$\llbracket \top \rrbracket_v = \llbracket \text{null} \rrbracket_v = \text{StructVal}, v, \emptyset$
$\text{nameof}(\text{string}) = \text{String}$ $\text{nameof}(\text{bool}) = \text{Boolean}$	$\text{nameof}(\text{number}) = \text{Number}$ $\text{nameof}(\text{collection}) = \text{List}$ $\text{nameof}(\text{rec-anon}) = \text{Record}$ $\text{nameof}(\text{rec-named } \nu) = \nu$

Figure 6. Type provider – generation of featherweight F# types from inferred structural types

The type provider generates the above type and the collection type is mapped to an F# type `list<PersonOrString>`. Given a structural document value s , the code to obtain the wrapped F# value is:

```
List.map (λx → new PersonOrString(x)) (getChildren(s))
```

The `PersonOrString` type contains one property for each of the union case. In the body, they check that the value is not `null` and that it has the right structure (using the `isTag` function). This checks that the value is a record named `Person` or a string, respectively. If the conditions are satisfied, the value is converted to the F# type corresponding to the case and wrapped in `Some`.

4.4 Inferring types from values

Before concluding the section on formalizing type providers, there is one more missing piece. The common supertype algorithm in Section 3 discusses the core of the type inference for structural values, but we have not yet clarified how exactly it is used. We address this in the present section.

Given a JSON, XML or CSV document, the F# Data implementation constructs a structural value s from the sample (this is straightforward, but Section 6 discusses some interesting aspects). The following defines a mapping $\llbracket - \rrbracket$ which turns a sample value s into a structural type σ :

$$\begin{aligned} \llbracket 0 \rrbracket &= \text{bit} & \llbracket 1 \rrbracket &= \text{bit} \\ \llbracket i \rrbracket &= \text{int} & \llbracket d \rrbracket &= \text{decimal} \\ \llbracket f \rrbracket &= \text{float} & \llbracket t \rrbracket &= \text{string} \\ \llbracket b \rrbracket &= \text{bool} & \llbracket \text{null} \rrbracket &= \text{null} \end{aligned}$$

$$\llbracket [s_1; \dots; s_n] \rrbracket = [\llbracket s_1 \rrbracket, \dots, \llbracket s_n \rrbracket]$$

$$\llbracket \nu_{\text{opt}} \{ \nu_1 \mapsto s_1, \dots, \nu_n \mapsto s_n \} \rrbracket =$$

$$\nu_{\text{opt}} \{ \nu_1 : \llbracket s_1 \rrbracket, \dots, \nu_n : \llbracket s_n \rrbracket \}$$

$$\llbracket s_1, \dots, s_n \rrbracket = \sigma_n \quad \text{where}$$

$$\sigma_0 = \top, \forall i \in \{1..n\}. \sigma_{i-1} \nabla \llbracket s_i \rrbracket \vdash \sigma_i$$

Primitive values are mapped to their corresponding types, with the exception of 0 and 1, which are inferred as `bit`. For records, we return a type with field types inferred based on individual values.

The interesting part is inferring type based on multiple samples. We overload the notation and write $\llbracket s_1, \dots, s_n \rrbracket$ for a type inferred from multiple samples. This uses the common supertype relation to find a common type for all values (starting with \top). This operation is used at the top-level (when calling type provider with multiple samples) and also when inferring the type of a collection.

5. Relativized type safety

Now that we have

Theorem 4 (Relativized safety). *Given samples s_1, \dots, s_n , inferred type $\sigma = \llbracket s_1, \dots, s_n \rrbracket$ and generated $\tau, e, L = \llbracket \sigma \rrbracket_x$ then:*

For all s' such that $\exists i. s_i :> s'$, then for any e such that $x : \text{StructVal} \vdash e : \tau$ it is the case that $e[x \leftarrow s'] \rightsquigarrow^ v$ and $\vdash v : \tau$.*

5.1 Discussion

alternative - check the whole file upfront based on the inferred type what if the sample is online??

6. Implementation

Turning JSON, CSV or XML to the idealized structured value presented here (esp. XML)

Error handling - solve issues just add another example

Unions become records with fields - so having null makes sense

Write API

JSON - inferring numbers from values (not just from actual value type - world bank example)

Heterogeneous collections

7. Conclusions

- Prime example of type providers, explores relativized type safety - Used by industry, downloads, tested, etc. - Predictable yet powerful type inference

A. Appendix Title

This is the text of the appendix, if you need one.

Acknowledgments

Acknowledgments, if needed.

References

[1] P. Q. Smith, and X. Y. Jones. ...reference text...